When Is Video Enough? Predicting Deformable Object Responses to New Interactions

Anonymous CVPR submission

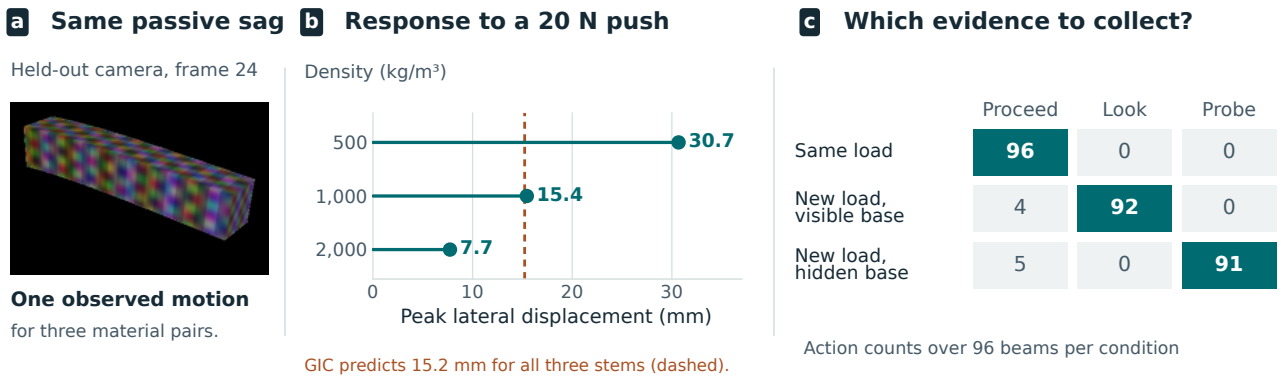


Figure 1. **The next interaction determines what the robot needs to know.** In the left and centre panels, three simulated stems have the same sag, yet move by 7.7–30.7 mm under a 20 N lateral push. GIC fitted to the passive images predicts 15.2 mm in all three. The right panel shows how, in a separate 96-case beam study, the initial tip observation can justify proceeding, observing the support, or probing, depending on the target and the accessible view.

Abstract

Interactive visual–physical world models aim to reconstruct scenes that can be both rendered from new viewpoints and simulated under new actions. Yet faithful reconstruction of an observed video need not identify the mechanics that determine a future intervention. Three simulated elastic stems exhibit indistinguishable gravity-driven motion but deflect by 7.7, 15.4, and 30.7 mm under the same applied force. We address this gap by inferring plausible material, support, and loading explanations and evaluating their action-conditioned *predictions rather than selecting a single visually convincing twin. For loading models affine in unknown force amplitudes, we analytically marginalize those amplitudes and retain their dependence on material and supports. For a proposed action, we propagate physical hypotheses through a forward solver, estimate the probability of exceeding a response-error tolerance, and weigh proceeding against another view or a calibrated probe. Held-camera rendering and held-out deformation experiments test visual–physical transfer, while separate spatial-inference and 96-case measurement studies examine task-dependent evidence acquisition. On a prospectively held-out diagonal stem push, one calibration interaction reduces 3D displacement error from 5.3–12.1 mm to 0.65–2.75 mm across three stems using RGB-D observations. The method proceeds for a predictable repeated action, generally requests a view when base motion is observable, and probes when the support is concealed. We also expose overconfidence under contact-model mismatch, delineating when task-specific risk estimates are trustworthy.*

1. Introduction

A useful *visual–physical world model* should do more than replay an observed video: it should reconstruct a scene that can be rendered from novel viewpoints *and* predict what happens when an embodied agent changes the scene. This joint ambition connects dynamic neural rendering to physical simulation and action-conditioned prediction. Dynamic Gaussian representations have made temporally coherent 3D reconstruction fast and photorealistic [26, 28]; physically grounded Gaussian and radiance-field systems bring explicit simulation into the representation [9, 15, 27]; and learned world models increasingly generate futures conditioned on manipulations or partial geometric controls [10, 12]. Together, these developments promise interactive digital twins for robotics and embodied reasoning.

However, *reconstructing observed motion and predict-*

ing a new physical response are different inverse problems. A method may explain the appearance, shape, and motion of the training sequence accurately while remaining wrong about the response to an unobserved force, contact, or support change. This is not necessarily an optimization or simulation defect. Several physically distinct worlds can generate the same observed trajectory, particularly when material stiffness, inertia, applied loads, and boundary compliance are incompletely known. No amount of photometric refinement of an intrinsically ambiguous trajectory can identify the missing mechanical scale. Accordingly, held-view PSNR and perceptual video quality are insufficient proxies for the accuracy of *counterfactual* force–displacement or impulse predictions. Recent physics-oriented benchmarks further motivate evaluating the physical state and response rather than rendered plausibility alone [5, 11].

The unresolved ambiguity is also *action dependent*. In Fig. 1, three simulated stems sag indistinguishably under gravity because their density and stiffness scale together. A known lateral push produces peak deflections from 7.7 to 30.7 mm, but repeating the original passive motion does not separate the worlds. The issue is therefore not whether every material parameter can be recovered uniquely. It is whether the explanations compatible with the available video predict sufficiently similar outcomes for *the action the robot intends to take*. Existing models already support novel interactions in important regimes [6, 9, 23, 25]; what remains insufficiently resolved is a prospective, tolerance-aware account of *when* an individual prediction is supported by its observations.

This viewpoint makes evidence acquisition part of physical world modeling. An additional camera may reveal out-of-plane deformation or a moving support, but it cannot necessarily recover an absolute stiffness scale from the same passive loading. A known-force probe may break that ambiguity, yet can be unnecessary when the intended action is already predictable. We seek to decide whether to *proceed, look, or probe* by comparing how each option is expected to change the error of the specified future response, rather than maximising visual reconstruction quality or parameter information independently of the task.

We propose a task-conditioned inference and decision layer for renderable physical twins. In the spatial identification setting, Gaussian fields parameterize uncertain material variation, support compliance, and load footprints in a finite-element model; Gaussian integration eliminates unknown load amplitudes where the observation map is affine in force. In known-command dynamic settings, calibrated image-derived surface motion weights candidate mechanical models, which are rolled out under a new force or contact. Both settings produce a distribution of possible *responses*, from which we estimate tolerance-exceedance risk and the value of an additional measurement. Our experiments separately test visual–physical transfer, joint physical inference, and cost-aware evidence selection; they do not claim an end-to-end photometric estimator for all scenes. Our contributions are (i) an action-specific ambiguity and risk formulation, (ii) joint material/support/load inference with task-driven view-or-probe selection, and (iii) prospective deformation, rendering, and measurement studies, including explicit controls for simulator mismatch.

2. Related Work

Dynamic 3D and 4D rendering. Deformable 3D Gaussians, 4D Gaussian Splatting, and Spacetime Gaussian Feature Splatting reconstruct time-varying scene appearance and support novel-view or novel-time synthesis [16, 26, 28]. These representations are powerful accounts of *where the scene appeared to move*; their standard deformation fields are indexed by time rather than a measured force or new robot command. *Limitation:* rendering an unseen viewpoint or interpolating a timestamp does not establish a physically valid response to an unseen intervention.

Physics-enabled rendering and forward simulation. PhysGaussian unifies 3D Gaussian rendering with material-point dynamics; Spring-Gaus couples Gaussian appearance to spring-mass systems; and PhysDreamer draws on video-generation priors to animate reconstructed assets under interactions [27, 30, 33]. PhysFlow, PhysSplat, GaussianFluent, and the learned Neural Gaussian Force Field extend simulation to richer materials, damage, or more efficient dynamics, while PhysGM predicts a simulation-ready Gaussian asset in a feed-forward pass [8, 14, 17, 18, 31]. Several methods genuinely accept physical controls or synthesize new interactions. *Limitation:* physical plausibility under a configured simulator or generative prior is not, by itself, calibrated evidence that its predicted response to a particular real-world load is correct.

Joint reconstruction and inverse physical identification. PAC-NeRF and GIC fit physical parameters through video, simulation, and rendering; MASIV learns a constitutive model from observed motion [3, 15, 32]. PhysTwin simulates changed grasp trajectories, M-PhyGs estimates spatially varying material for transfer between manipulations, and ReconPhys predicts monocular-video-derived appearance and physical attributes [9, 23, 24]. Most directly, PhysConvex uses deformable convexes as common rendering and continuum-simulation supports, jointly refines physical parameters from multiview videos, quantitatively evaluates eight held-out continuation frames and illustrates modified forces and boundaries [25]. Its published new-force and new-boundary demonstrations should not be dismissed as absent, but neither should they be equated with prospectively calibrated force-response risk across ambiguous physical worlds.

Limitation: these methods provide capable fitted physical twins, but their reported evaluations do not directly test whether a particular next response remains non-identifiable under several observationally compatible material, load, and support configurations.

Action-conditioned and generative world models. P3Sim conditions a probabilistic RGB/depth/flow model on partial geometric motion signals derived from specified camera or object transforms; its quantitative benchmarks address novel-view synthesis and image-based 3D editing, whereas its deformable manipulation and collision demonstrations are primarily qualitative [12]. Physical Object Understanding learns a physically controllable probabilistic object model [22]. Dexterous World Models predicts hand-motion-conditioned video; Goal Force explicitly conditions on force-related goals, PhyCo conditions on physical properties, and video–action models condition on robot actions [7, 10, 19, 34]. Thus, *missing action conditioning* is not the gap. *Limitation:* geometric edits, generated futures, and action-progress scores are not evaluations of calibrated millimetre-level deformation or Newton-valued reaction-force error across observationally ambiguous mechanical explanations.

Interactive calibration, uncertainty, and active measurements. Latent-force Gaussian processes and Bayesian experimental design provide ways to infer unmeasured forcing and choose informative observations [1, 2, 20, 21]. KnockGS calibrates two global effective elasticity/density scales from a known Probe-A particle response, freezes its response-feature regression estimate and tests unseen probes of different direction and magnitude [6]. This is a direct and important probe-to-probe prediction precedent, though its object-specific candidate-response library and partial released code do not establish a general end-to-end image-to-simulator pipeline. PhysCoRe infers spatial material from observed RGB-D motion and corrects an MPM solver with a learned velocity residual; it reports real-object dynamics prediction and visualises material-confidence changes as a robot probes objects [29]. The published confidence-guided exploration section does not report a matched-budget quantitative policy comparison that prospectively selects among view, probe and proceed. *Limitation:* both works address important components of interaction-informed prediction, but neither explicitly minimises tolerance-exceedance risk for the *specified* next action over alternative evidence sources and their acquisition costs.

Physical-fidelity datasets and evaluations. PhysGaia evaluates multi-body physical consistency in dynamic novel-view synthesis, PhysInOne provides broad annotated physical phenomena, and Deform360 combines multiview visuotactile evidence for deformable-world modeling [11, 13, 35]. GaussianBench further separates physical failures, unsupported capabilities, and invalid tests under interventions [5]. *Limitation:* broad reconstruction or physics-fidelity evaluations do not themselves answer, prospectively, whether the currently observed object is predictable *for one specified next action*, nor whether looking or probing is the cheaper way to make that prediction reliable.

3. Predicting an Interaction from Partial Evidence

Let D contain the available motion observations and known input commands. A physical explanation z consists of the relevant material, boundary, and loading variables. A forward model produces the visible observations $g(z)$ and the response $q(z, a_\star)$ to a specified future interaction $a_\star$. The target may be a trajectory of surface points or an impulse delivered during contact. The command is available at prediction time. Its realized response is withheld.

When does an ambiguity matter? Suppose two explanations z_1, z_2 produce the same noiseless observations. Every estimator using those observations and the same target command must give them the same point prediction $\hat{q}$. For a response metric d the triangle inequality gives

$$\max_{i=1,2} d\big(q(z_i, a_\star), \hat{q}\big) \;\geq\; \tfrac{1}{2} d\big(q(z_1, a_\star), q(z_2, a_\star)\big). \qquad (1)$$

If the two responses differ by more than 2δ, no single estimate can meet tolerance δ in both worlds. In Fig. 1, even a perfect passive fit must miss at least one extreme stem by about 11.5 mm in peak lateral deflection. Changing the task to a repeated gravity release collapses this separation. This bound is a worst-case statement. Estimating the probability of a threshold violation additionally requires a distribution over the plausible explanations.

Prediction and a decision to measure. We assign posterior weights to the physical hypotheses using D. Under squared error the point forecast is the posterior mean, $\hat{q} = \mathbb{E}[q(z, a_\star) \mid D]$. For any specified point forecast, its estimated task risk is

$$r(D, a_\star) = \Pr\big[d\big(q(z, a_\star), \hat{q}(D, a_\star)\big) > \delta \mid D\big]. \qquad (2)$$

For a finite posterior this is the sum of the weights of hypotheses outside the tolerance. It can be large even for an excellent image fit. Its calibration depends on the forward models and sensor noise represented in the posterior.

Let s be a candidate measurement, y_s its possible outcome and c_s its cost in units of threshold failure. We choose

the lowest-cost option from

$$J(\text{proceed}) = r(D, a_\star), \tag{3}$$

$$J(s) = c_s + \mathbb{E}_{y_s|D,s}\big[r(D \cup \{y_s\}, a_\star)\big]. \tag{4}$$

For each hypothetical observation we update *both* the weights and the point prediction. The expectation uses possible measurement outcomes available from the model, never the true future response. The beam experiment uses threshold loss and different costs for looking and pushing. The spatial field experiment evaluates an analogous squared-error design for a fixed number of scalar measurements.

4. Physical Explanations and Observation Models

Spatial material, support and load. Figure 2 shows the unknown-load case. In the joint-field study the object is a known, small-strain elastic volume with displacement u_t from its rest configuration. Its positive Young's modulus is $E(X) = \exp(m_0 + m_1\chi(X))$ at material coordinate X, where χ is the volume-centred difference of two normalized Gaussian features. Two additional Gaussian surface footprints describe a compliant support and an applied traction. Each surface footprint integrates to one over its surface. Its position describes *where* the interaction acts, while each traction amplitude gives a total force in newtons. The support has stiffness k_b and energy $\frac{1}{2}k_b \int_{\Gamma_b} w_b(X)\|u\|^2 dA$. Two signed traction modes point downward on the top and laterally on a side. They share an inferred longitudinal centre c_f. The field hypothesis is $\theta = (m_0, m_1, c_b, \log k_b, c_f)$.

For each candidate, the tetrahedral stiffness operator K_θ and the unit-load map B_θ give

$$K_\theta u_t = B_\theta a_t + f_t^{\text{known}}, \qquad y = H_\theta a + d_\theta + \epsilon. \tag{5}$$

Here a stacks two force amplitudes at each of eight times. Solves for unit loads and projection into calibrated cameras define H_θ; d_θ is the known-load contribution. Scaling material stiffness, support stiffness and an unknown force by the same factor preserves the unknown-load displacement. An independently calibrated force breaks this symmetry. A Gaussian support footprint locates a mechanical interaction; a rendering Gaussian carries appearance. Only the former enters the stiffness operator.

Marginalizing unknown load histories. Conditional on θ, let $a \sim \mathcal{N}(0, \Lambda)$ with temporally correlated amplitudes and $\epsilon \sim \mathcal{N}(0, R)$. Gaussian integration gives

$$C_\theta = R + H_\theta \Lambda H_\theta^\top, \tag{6}$$

$$p(y \mid \theta) = \mathcal{N}(y; d_\theta, C_\theta), \tag{7}$$

$$\mu_\theta = \Lambda H_\theta^\top C_\theta^{-1}(y - d_\theta), \tag{8}$$

$$S_\theta = \Lambda - \Lambda H_\theta^\top C_\theta^{-1} H_\theta \Lambda. \tag{9}$$

The last two quantities describe the conditional force posterior. The likelihood includes $\log|C_\theta|$ as well as the residual. Thus a stiff body can fit the same measured displacement as a soft body while requiring a different force distribution. Predicting a hidden displacement under the original load uses the corresponding force conditional; a known reset load requires a forward solve for each material/support hypothesis. In both cases we integrate over the posterior weights. The unknown-force calculation applies where the observation model is affine in force amplitudes. Contact and impact forecasts below use known commands and direct rollouts instead.

An enumerated 5^5 dictionary provides exact weights for its 3,125 spatial hypotheses. An independent implementation uses adaptive tempered sequential Monte Carlo in the five continuous coordinates [4], retaining the Gaussian force conditional. The continuous study replays the finite planner's actions. A hierarchical variant additionally infers one of five force amplitude scales under a Matérn-$3/2$ time prior. This accounts for a force prior itself supplying apparent information about stiffness. Discretization, priors, and sensitivity to temporal mismatch are given in the supplement.

A view and a calibrated reset. For the spatial field study the target concatenates hidden displacements under the original loads and responses to four *new* known loads, with 192 scalar coordinates in each group. The design minimizes expected coordinate variance averaged over this fixed target. A new projected track on the original episode updates both the physical weights and conditional forces. A reset probe supplies a response to a known load and updates the physical weights alone. We allow four scalar measurements and approximate each one-step predictive expectation with 128 common-random-number draws. For the separate beam decision study, bending and base compliance replace the five spatial coordinates. A single tip observation constrains their sum, while either a view of the moving support or a calibrated push at another location can distinguish their contributions.

From video to dynamic response. The dynamic trials take calibrated cameras, a metric rest shape, and known load commands. The completed RGB-D stem study follows surface texture with pyramidal Lucas–Kanade tracking initialized from known rest points; noisy depth supplies the third coordinate. The RGB-only protocol instead follows each point in two overlapping fitting cameras and triangulates its position. Geometry and rest-point initialization are additional inputs to our estimator. GIC and PAC-NeRF reconstruct their object from video. The reserved camera participates only in evaluation.

For the stem, each candidate specifies E/ρ, ρ and Poisson's ratio; the support is clamped. A bank of coarse linear

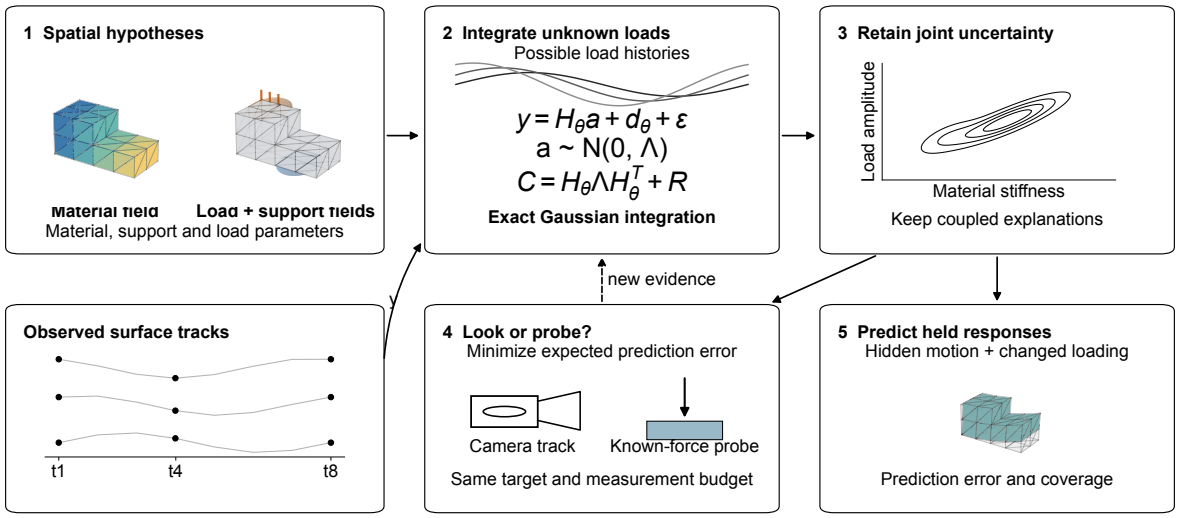


Figure 2. **The unknown-load branch of the method.** Each spatial hypothesis defines a response to unit forces. Integrating over force amplitudes assigns weights to material and support hypotheses and preserves their conditional force distributions. Candidate views and probes are compared in terms of the resulting change in predictive uncertainty. In the known-command dynamic experiments we use the measured command directly.

elastodynamic responses predicts motion under gravity and an applied force. With observed marker displacements y_i, predicted displacements $g_i(\theta)$ and observation noise σ, the weights after each episode satisfy

$$w_\theta \propto p(\theta) \exp\left[-\frac{1}{2\sigma^2}\sum_i \|y_i - g_i(\theta)\|^2\right]. \quad (10)$$

We keep the noise estimate from the initial fit for the calibration update. We apply the held-out command to each weighted hypothesis and compute its predictive mean and threshold risk. The RGB-only block protocol uses an analogous bank over inclusion position, stiffness contrast, background modulus and Poisson ratio, with a homogeneous block also permitted. Both published competitors use a single global modulus in the configured benchmark; a test of inference alone would require the same heterogeneous material family for all estimators. The striker protocol propagates the calibrated stem posterior through a compliant contact law and scores momentum transfer from the freely evolving sphere. The calibration command is fixed in these three scenarios. Selective acquisition is evaluated separately.

5. Experiments and Results

5.1. Does the observed motion identify the next response?

Fit and forecast on the same camera. We fitted the released GIC [3] and PAC-NeRF [15] implementations to eight frames of the synthetic Torus sequence from ten cameras. We reserved camera 0 throughout and measured mean per-frame foreground PSNR there. GIC's deformation network reproduced the observed motion at 20.47 dB, ahead of its physics replay at 17.79 dB. Over the next six frames the deformation network scored 8.29 dB, below its frozen last frame (9.50 dB), while the physics forecast reached 18.44 dB. PAC-NeRF's physics forecast scored 18.01 dB. This experiment tests continuation of observed motion. A separate intervention on three released ReconPhys assets [24] probes absolute scale. Doubling mass, stiffness and damping together gave 94.6–100.1 dB between simulated drop movies, yet the two versions responded differently to the same uniform force per anchor.

A calibrated intervention on visually identical stems. Three stems have densities of 500, 1000 and 2000 kg/m^3 and proportional elastic moduli of 2, 4 and 8 MPa. They sag identically under gravity. A 20 N lateral load at three quarters of the free length produces peak lateral displacements of 30.7, 15.4 and 7.7 mm in the nonlinear FEM reference. To test transfer, we sealed predictions of a new 25 N diagonal load at the same location before generating its reference. All methods received the known commands. GIC reconstructed the stem from ten RGB views and fixed density at 1000 kg/m^3. Our estimator used 16 surface points tracked in two RGB views with noisy rendered depth and known rest geometry. Its coarse linear tetrahedral model shares its element family with the nonlinear reference.

Table 1. New diagonal push. RMS 3D displacement error (mm) across four stations and 37 frames, before and after one 40 N midpoint poke.

Density (kg/m^3)	500	1000	2000
GIC, RGB, before	14.33	3.08	7.18
GIC, RGB, after	12.98	1.88	8.08
Ours, RGB-D, before	8.74	5.30	12.08
Ours, RGB-D, after	2.75	1.12	0.65

Calibration brought our errors below the 3 mm tolerance in all three cases (Table 1). In the earlier 20 N lateral-push trial, the same-bank posterior mode and posterior mean gave the same calibrated point forecast. GIC improved on the two lighter stems and crossed the threshold at its assumed density. On the withheld camera its deformation network rendered passive sag at 24.9 dB and its physics replay at 20.6 dB. For the light stem, the passive physics forecast scored 16.3 dB on the diagonal push. A subsequent rendering analysis (Fig. 3) drove the same fixed GIC appearance with our calibrated motion or the reference motion, yielding 21.27 and 21.49 dB, respectively. We sealed the motion forecasts prospectively and finalized the common-renderer analysis and its foreground-error aggregation separately.

Choosing whether to measure. In an independent beam study, 96 draws of bending and base compliance each faced three task or visibility conditions. All policies saw the same noisy tip measurement. With a tolerance of 0.50 mm, a view and a calibrated push cost 0.015 and 0.040 in failure-probability units. For a repeated tip load the risk policy proceeded in all 96 cases (0.104 mm RMSE). For a shifted load the initial RMSE was 0.939 mm. With a visible base it acquired a view in 92 cases (0.073 mm). With a concealed base it probed in 91 (0.119 mm). It proceeded in the remaining cases. Against a cost-aware rule using Gaussian moments, the paired decision-loss difference (full minus Gaussian) was -0.00033 with a 95% world-bootstrap interval of $[-0.00085, 0]$.

Image and impulse rankings under contact. In an earlier six-family study, adapted Spring-Gaus [33] and PhysTwin [9] mechanics received identical calibration tracks, geometry, density and boundary conditions. A common renderer displayed their forecasts. On eight held-out striker impacts across four material worlds, Spring-Gaus exceeded PhysTwin's foreground PSNR in every world (mean advantage 1.22 dB), yet its 3D motion error was higher (4.80 versus 4.01 mm) and its net impulse error was 0.0419 versus 0.0231 N s. On eight curved-obstacle pushes Spring-Gaus gained 2.35 dB and reduced motion error from 4.97 to 3.15 mm. These measurements characterize the two adapted mechanics implementations under supplied tracks.

Table 2. **Prediction after four scalar observations.** Means over 128 matched cases. Hidden and new-load errors are coordinate RMSE in mm. Material error is relative; force error is in newtons. Fixed denotes the development-selected mixed schedule.

Policy	Hidden mm	New load mm	Material rel.	Force N
Initial	0.076	0.024	0.451	0.042
Active	0.045	0.014	0.347	0.030
Fixed	0.057	0.012	0.322	0.033
Fixed View	0.050	0.022	0.421	0.032
Fixed Probe	0.074	0.010	0.277	0.040
Random	0.059	0.014	0.343	0.034

Sensitivity to model error. A deformable body in a concave bowl was pushed or tilted once before a stronger transverse push. Across 16 material/contact settings with two noise draws each, the coarse model predicted 1.86% failure after adaptive calibration, compared with 43.75% realized against a finer simulation. A mixture including the finer solver by construction improved coverage. Risk estimates consequently depend on which mechanical discrepancies the posterior represents.

Three-scenario comparison protocol. The RGB-only protocol pairs GIC, PAC-NeRF and our estimator on a new stem push, an off-centre press on a block with a hidden stiff region, and a freely rebounding striker. The push and impact reuse the stem geometry. The protocol specifies triangulated RGB observations and held-out displacement or impulse. Aggregate results for this comparison are pending.

5.2. Joint spatial identification under unknown loading

We also evaluate when a view or a known force provides better evidence in the presence of unknown loading. We use a block and stepped block of bounding dimensions $120 \times 60 \times 60$ mm. Inference starts from 24 scalar observations of projected motion on the visible surface: four material points, two orthographic components and three loading times. Additional tracks use an oblique camera; a geometric visibility check covers all deformed test states. Sensor noise is 0.02 mm. The force standard deviation and calibrated probe amplitude are 0.08 N. The two geometry-specific development sets select fixed schedules before held-out evaluation. The supplement describes the discretization and all independent cohorts.

Acquisition and the target. Across 128 matched cases, four adaptive measurements reduce combined coordinate

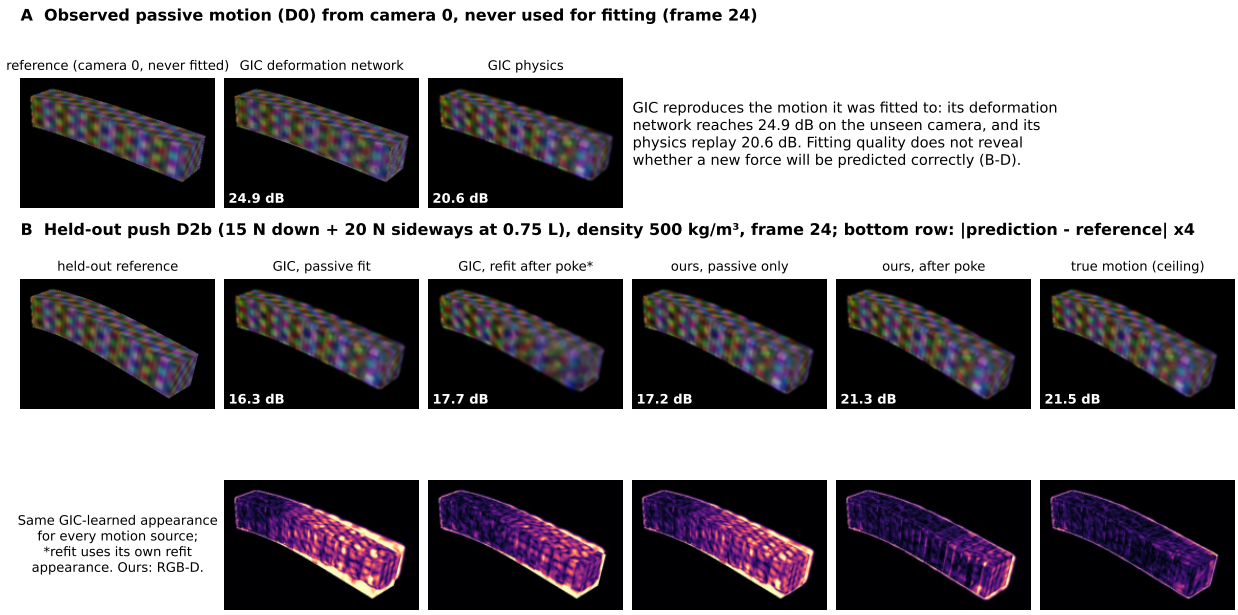


Figure 3. **Rendering the measured and predicted stem motion from a withheld camera.** The upper row shows a passive frame. The lower panels show the new diagonal push on the light stem with pixel differences magnified fourfold. Numbers in the images are sequence foreground PSNR. Comparisons primarily use one frozen GIC appearance; the GIC refit uses its own appearance. The motion predictions were sealed before reference generation; the common-appearance analysis was finalized subsequently.

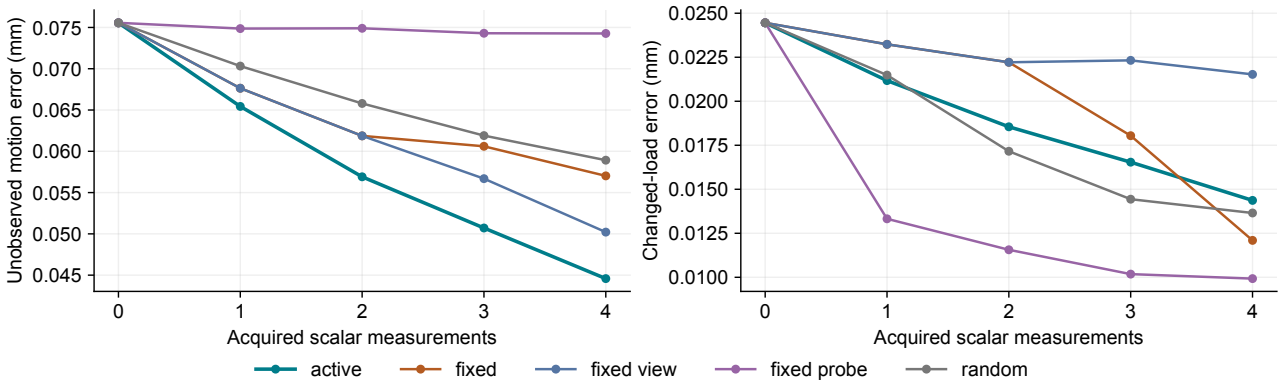


Figure 4. **Measurement choice depends on the target.** An additional camera can expose the original hidden motion, whereas a known probe directly constrains response under a changed load. Curves average all matched surface-track cases.

RMSE from 0.0578 to 0.0340 mm. Fixed camera-only, fixed mixed and random policies reach 0.0403, 0.0423 and 0.0438 mm. Against the fixed mixed schedule, the paired difference is -0.0083 mm (95% case-bootstrap interval $[-0.0110, -0.0058]$). The adaptive policy spends 439 of its 512 measurements on camera observations. Probes alone give the best prediction under known new loads, at 0.0099 mm against 0.0144 mm for adaptive selection, but leave more uncertainty in the original hidden motion (Table 2, Fig. 4). In another 128 cases with uncertain force scale, four fixed views reduce hidden-motion error by 33.4%, while four calibrated probes reduce new-load error by 48.4%. Both branches start with the same evidence.

What retaining dependence changes. Retaining the conditional force distribution for each material/support hypothesis gives a mean prediction-interval width of 0.0717 mm; replacing those conditionals with a common moment-matched distribution gives 0.0863 mm. Coverage is 92.0% versus 93.9% at nominal 90%, and hidden-motion errors are nearly equal (0.0446 and 0.0449 mm). In this experiment the practical gain is narrower intervals at similar accuracy. On a

separate continuous-prior cohort, sequential Monte Carlo gives 0.0294 mm after the recorded adaptive actions versus 0.0350 mm after fixed mixed observations.

Calibration outside the fitted family. Finite-model prediction intervals cover 92.0% of matched target coordinates. Coverage falls to 86.4% for continuous parameters, 86.1% when the generator uses a finer mesh, and 79.2% for a shorter correlation length in the true force history. A follow-up that mixed the temporal length and amplitude on 64 fresh stress cases improved force coverage from 72.4% to 83.9%; inferred material error increased. More observations can narrow a posterior around the wrong explanation.

6. Discussion

Visual evidence is informative about the motion it reveals. A change in load, support or contact tests which mechanical explanations remain compatible with the response a robot needs. The observed rendering scores, the calibrated-force trials and the choice to look or probe address different aspects of that question. A camera change can disclose support movement; even many extra views of an identical 3D sag leave the force scale ambiguous. For the diagonal push, the GIC simulator supplied with the true material parameters still has 7.1–9.3 mm station error. Its reconstructed shape, fixed-density identification and MPM dynamics all affect the comparison. Our estimator knows the rest shape, receives depth in the completed prospective trial, and uses the same element family as the finite simulation reference. These numbers establish a controlled counterexample and a repeatable calibration effect under the stated sensing conditions. Ranking full systems on real RGB footage requires a matched-input test.

Accurate risk predictions require a posterior broad enough to represent sensor and simulator error. The bowl result quantifies the cost of ignoring this condition. Similarly, the beam policy's small advantage over a cost-aware Gaussian approximation has an interval including zero. A real-object transfer experiment would need measured inputs for calibration, an independently observed response to a different interaction, and evaluation of failures across objects. Known geometry, track initialization and the single-step cost model remain practical restrictions. In the present RGB protocol, the physical likelihood uses tracked motion; image rendering provides an independent visual readout. Weighting hypotheses directly by visible silhouettes or reliable multiview pixel residuals is a concrete extension to test on fresh episodes.

7. Conclusion

An observation identifies a response only to the extent that compatible models agree about the intended interaction. Propagating joint physical uncertainty through that interaction gives a testable prediction and a reason to seek specific evidence. Our experiments show when another view or a known probe resolves the ambiguity, and how a misspecified simulator can make that decision overconfident.

LLM Usage Statement

OpenAI ChatGPT and coding agents, and Anthropic Claude, assisted with implementation, synthetic experiment orchestration, numerical and plotting scripts, literature organization, and manuscript preparation. ChatGPT was also used for reviewer-style feedback, exploration of experimental hypotheses, and editing and drafting text. Numerical figures were produced by scripts from simulation outputs or recorded scores. The authors take responsibility for the submitted text, references, code, and results.

References

[1] Mauricio Álvarez, David Luengo, and Neil D. Lawrence. Latent force models. In *Proceedings of the Twelfth International Conference on Artificial Intelligence and Statistics*, pages 9–16, 2009. 3

[2] Ahmed Attia, Alen Alexanderian, and Arvind K. Saibaba. Goal-oriented optimal design of experiments for large-scale bayesian linear inverse problems. *Inverse Problems*, 2018. 3

[3] Junhao Cai, Yuji Yang, Weihao Yuan, Yisheng He, Zilong Dong, Liefeng Bo, Hui Cheng, and Qifeng Chen. GIC: Gaussian-informed continuum for physical property identification and simulation. In *Advances in Neural Information Processing Systems*, 2024. 2, 5

[4] Pierre Del Moral, Arnaud Doucet, and Ajay Jasra. Sequential monte carlo samplers. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 68(3):411–436, 2006. 4

[5] Chukwudalu Dumebi-Kachikwu. GaussianBench: Physics-fidelity evaluation for gaussian scene representations, 2026. 2, 3

[6] Chenchen Ge, Hanwen Shen, Bowen Jing, Jiyuan Cai, Xiaofeng Wang, Hongsen Lei, Weitao Zhou, Dandan Zhang, and Haibao Yu. KnockGS: Interaction-grounded calibration of physical gaussian representations. *arXiv preprint arXiv:2608.27365*, 2026. 2, 3

[7] Nate Gillman, Yinghua Zhou, Zitian Tang, Evan Luo, Arjan Chakravarthy, Daksh Aggarwal, Michael Freeman, and Chen Sun. Goal force: Teaching video models to accomplish physics-conditioned goals. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 20077–20087, 2026. 3

[8] Bei Huang, Yixin Chen, Ruijie Lu, Gang Zeng, Hongbin Zha, Yuru Pei, and Siyuan Huang. GaussianFluent: Gaussian simulation for dynamic scenes with mixed materials. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 21583–21593, 2026. 2

[9] Hanxiao Jiang, Hao-Yu Hsu, Kaifeng Zhang, Hsin-Ni Yu, Shenlong Wang, and Yunzhu Li. PhysTwin: Physics-informed reconstruction and simulation of deformable objects from

videos. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2025. 1, 2, 6

[10] Byungjun Kim, Taeksoo Kim, Junyoung Lee, and Hanbyul Joo. Dexterous world models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 29663–29673, 2026. 1, 3

[11] Mijeong Kim, Gunhee Kim, Jungyoon Choi, Wonjae Roh, and Bohyung Han. PhysGaia: A physics-aware benchmark with multi-body interactions for dynamic novel view synthesis. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 22604–22614, 2026. 2, 3

[12] Wanhee Lee, Klemen Kotar, Rahul Mysore Venkatesh, Jared Watrous, and Daniel L. K. Yamins. Perceptual 3D simulation with physical world modeling. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 27219–27228, 2026. 1, 3

[13] Hongyu Li, Wanjia Fu, Xiaoyan Cong, Zekun Li, Binghao Huang, Hanxiao Jiang, Xintong He, Yiqing Liang, Rao Fu, Tao Lu, Srinath Sridhar, Kevin A. Smith, George Konidaris, and Yunzhu Li. Deform360: A massive multi-view visuotactile dataset for deformable world models. In *European Conference on Computer Vision*, 2026. 3

[14] Shiqian Li, Ruihong Shen, Junfeng Ni, Chang Pan, Chi Zhang, and Yixin Zhu. Learning physics-grounded 4D dynamics with neural gaussian force fields. In *International Conference on Learning Representations*, 2026. 2

[15] Xuan Li, Yi-Ling Qiao, Peter Yichen Chen, Krishna Murthy Jatavallabhula, Ming Lin, Chenfanfu Jiang, and Chuang Gan. PAC-NeRF: Physics augmented continuum neural radiance fields for geometry-agnostic system identification. In *International Conference on Learning Representations*, 2023. 1, 2, 5

[16] Zhan Li, Zhang Chen, Zhong Li, and Yi Xu. Spacetime gaussian feature splatting for real-time dynamic view synthesis. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 8508–8520, 2024. 2

[17] Zhuoman Liu, Weicai Ye, Yan Luximon, Pengfei Wan, and Di Zhang. Unleashing the potential of multi-modal foundation models and video diffusion for 4D dynamic physical scene simulation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 11016–11025, 2025. 2

[18] Chunji Lv, Zequn Chen, Donglin Di, Weinan Zhang, Hao Li, Wei Chen, Yinjie Lei, and Changsheng Li. PhysGM: Large physical gaussian model for feed-forward 4D synthesis. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2026. 2

[19] Sriram Narayanan, Ziyu Jiang, Srinivasa Narasimhan, and Manmohan Chandraker. PhyCo: Learning controllable physical priors for generative motion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 41892–41902, 2026. 3

[20] Rajdip Nayek, Souvik Chakraborty, and Sriram Narasimhan. A gaussian process latent force model for joint input-state estimation in linear structural systems. *Mechanical Systems and Signal Processing*, 2019. 3

[21] T. J. Rogers, K. Worden, and E. J. Cross. On the application of gaussian process latent force models for joint input-state-parameter estimation: With a view to bayesian operational identification. *Mechanical Systems and Signal Processing*, 140:106580, 2020. 3

[22] Rahul Venkatesh, Klemen Kotar, Lilian Naing Chen, Wanhee Lee, Gia Ancone, Seungwoo Kim, Luca Thomas Wheeler, Jared Watrous, Honglin Chen, Daniel Bear, Stefan Stojanov, and Daniel L. K. Yamins. Physical object understanding with a physically controllable world model. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 2593–2602, 2026. 3

[23] Norika Wada, Kohei Yamashita, Ryo Kawahara, and Ko Nishino. M-PhyGs: Multi-material object dynamics from video. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Findings*, 2026. 2

[24] Boyuan Wang, Xiaofeng Wang, Yongkang Li, Zheng Zhu, Yifan Chang, Angen Ye, Guosheng Zhao, Chaojun Ni, Guan Huang, Yijie Ren, Yueqi Duan, and Xingang Wang. ReconPhys: Reconstruct appearance and physical attributes from single video. In *European Conference on Computer Vision*, pages 492–509, 2026. 2, 5

[25] Dan Wang, Xinrui Cui, Serge Belongie, and Ravi Ramamoorthi. PhysConvex: Physics-informed dynamic convex fields for reconstruction and simulation. In *European Conference on Computer Vision*, 2026. 2

[26] Guanjun Wu, Taoran Yi, Jiemin Fang, Lingxi Xie, Xiaopeng Zhang, Wei Wei, Wenyu Liu, Qi Tian, and Xinggang Wang. 4D gaussian splatting for real-time dynamic scene rendering. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 20310–20320, 2024. 1, 2

[27] Tianyi Xie, Zeshun Zong, Yuxing Qiu, Xuan Li, Yutao Feng, Yin Yang, and Chenfanfu Jiang. PhysGaussian: Physics-integrated 3D gaussians for generative dynamics. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024. 1, 2

[28] Ziyi Yang, Xinyu Gao, Wen Zhou, Shaohui Jiao, Yuqing Zhang, and Xiaogang Jin. Deformable 3D gaussians for high-fidelity monocular dynamic scene reconstruction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 20331–20341, 2024. 1, 2

[29] Haocheng Yin, Shuohan Tao, Yongsheng Chen, and Lu Gan. PhysCoRe: Physics-corrected residual world models for material-aware deformable dynamics. In *Conference on Robot Learning*, 2026. 3

[30] Tianyuan Zhang, Hong-Xing Yu, Rundi Wu, Brandon Y. Feng, Changxi Zheng, Noah Snavely, Jiajun Wu, and William T. Freeman. PhysDreamer: Physics-based interaction with 3D objects via video generation. In *European Conference on Computer Vision*, 2024. 2

[31] Haoyu Zhao, Hao Wang, Xingyue Zhao, Hao Fei, Hongqiu Wang, Chengjiang Long, and Hua Zou. PhysSplat: Efficient physics simulation for 3D scenes via MLLM-guided gaussian splatting. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 5242–5252, 2025. 2

[32] Yizhou Zhao, Haoyu Chen, Chunjiang Liu, Zhenyang Li, Charles Herrmann, Junhwa Hur, Yinxiao Li, Ming-Hsuan Yang, Bhiksha Raj, and Min Xu. Toward material-agnostic system identification from videos. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2025. 2

[33] Licheng Zhong, Hong-Xing Yu, Jiajun Wu, and Yunzhu Li. Reconstruction and simulation of elastic objects with spring-mass 3D gaussians. In *European Conference on Computer Vision*, 2024. 2, 6

[34] Pengfei Zhou, Shengcong Chen, Di Chen, Jiaxu Wang, Rongjun Jin, Bingwen Zhu, Yike Pan, Songen Gu, et al. tau0-WM: A unified video-action world model for robotic manipulation, 2026. 3

[35] Siyuan Zhou, Hejun Wang, Hu Cheng, Jinxi Li, Dongsheng Wang, Junwei Jiang, Yixiao Jin, Jiayue Huang, et al. PhysI-nOne: Visual physics learning and reasoning in one suite. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 33131–33142, 2026. 3